

Write Up

By Luca Comba and Pam Savira.

Source code available at: https://github.com/lukfd/mlops_final

Introduction

Building off of https://github.com/lukfd/lab4_model_training we decided to add Kubernetes in Docker to allow us to have a SQS consumer to run inferences.

Steps

This setup runs LocalStack, Airflow, FastAPI, and a nested Kind cluster inside Docker Compose. Airflow talks to the Kind API through a shared kubeconfig, and the consumers inside Kind reach LocalStack via a Kubernetes ExternalName service.

Docker setup

1. Ensure Docker has at least 8 GB of RAM allocated.
2. Make the LocalStack init script executable (only once):

```
chmod +x scripts/localstack_init.sh
```

3. Build and start the stack:

```
docker compose up -d --build
```

4. Confirm the Kind cluster from Airflow:

```
make kind-get-nodes
```

Note: the Kind API server is pinned to port 6443 using the project `kind-config.yaml`. The provisioner rewrites the kubeconfig to `https://docker-host:6443` so Airflow can reach the nested cluster over the Compose network.

If `make kind-get-nodes` fails, verify the control-plane port mapping inside the DinD host:

```
docker compose exec -T docker-host docker ps
```

You should see the Kind control plane bound to `0.0.0.0:6443->6443/tcp`.

Train

Run the cancer training pipeline in Airflow so it uploads the promoted model to S3 (LocalStack):

- Trigger the `ml_training_pipeline_v2` DAG in the Airflow UI.

This produces: - models/cancer/<version>/cancer.pkl - models/cancer/<version>/metrics.json
- models/cancer/<version>/metadata.json - models/cancer/promoted_model.json

SQS and consumer

1. Build the consumer image and load it into the Kind cluster:

```
make kind-build-consumer
make kind-load-consumer
## Or one shot:
## make kind-refresh-consumer
```

2. Apply the Kubernetes deployment:

```
make kind-apply-consumer
```

3. (Optional) Scale the consumer replicas:

```
make kind-scale-consumer
```

4. Trigger the `sqs_populate_inference_queue` DAG in Airflow. It loads the breast cancer test split and sends one message per record to SQS with `record_id` and `features`.

As the consumers poll SQS, they load the promoted model from S3 on startup, perform inference, and write each result to `s3://lab4-model-training/predictions/<record_id>.json`.

Results

We successfully ran the entire pipeline locally, and we are sure that we can port these results into AWS or any other cloud provider.

Like described in the previous “Steps” section, we ran the `make up`. We then ran the `ml_training_pipeline_v2` DAG in Airflow.

Then we created all resources in kubernetes

```
@mlops_final $ make kind-build-consumer
docker compose -f docker-compose.yml exec -T docker-host \
    docker build -t inference-consumer:latest /opt/project/consumer
#0 building with "default" instance using docker driver

#1 [internal] load build definition from Dockerfile
#1 transferring dockerfile: 251B done
#1 DONE 0.0s

#2 [internal] load .dockerignore
#2 transferring context: 2B done
#2 DONE 0.0s
```

```
#3 [internal] load metadata for docker.io/library/python:3.11-slim
#3 DONE 1.0s
```

```
#4 [internal] load build context
#4 transferring context: 3.44kB done
#4 DONE 0.0s
```

```
#5 [1/5] FROM docker.io/library/python:3.11-slim@sha256:efde372d959bc3cb6848a31658b603877b99
```

... truncated logs ...

```
@mlops_final $ make kind-apply-consumer
docker compose -f docker-compose.yml exec airflow-webserver \
    kubectl --kubeconfig /opt/project/kubeconfig.yaml apply -f /opt/project/k8s/
deployment.apps/inference-consumer created
@mlops_final $ make kind-scale-consumer
docker compose -f docker-compose.yml exec airflow-webserver \
    kubectl --kubeconfig /opt/project/kubeconfig.yaml scale deployment inference-
deployment.apps/inference-consumer scaled
```

We made sure that the kubernetes resources were up and running

```
@mlops_final $ docker compose exec airflow-webserver kubectl --kubeconfig /opt/project/kubeconfig.yaml get pods
NAME                                READY   STATUS    RESTARTS   AGE
coredns-6f6b679f8f-ldz9s           1/1     Running   0           20m
coredns-6f6b679f8f-sv2pk           1/1     Running   0           20m
@mlops_final $ docker compose exec airflow-webserver kubectl --kubeconfig /opt/project/kubeconfig.yaml get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
localstack ExternalName <none>        localstack.ai-pipeline-net <none>      21m
```

Then we ran the `sqs_populate_inference_queue` DAG which kicked off the inferences by sending new messages to SQS so that the kubernetes consumer could pick them up.

```
docker compose exec airflow-webserver kubectl --kubeconfig /opt/project/kubeconfig.yaml logs -f inference-consumer
```

No real messages were created in the pod, so we then created a POST `/predictions` api endpoint to retrieve the predictions generated that were uploaded to S3.

```
{
  "count": 10,
  "items": [
    {
      "record_id": "sample_0016",
      "prediction": 1,
      "timestamp": "2026-05-09T05:03:15.001289+00:00"
    },
    {

```

```

        "record_id": "sample_0019",
        "prediction": 0,
        "timestamp": "2026-05-09T05:03:15.001889+00:00"
    },
    {
        "record_id": "sample_0020",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.006647+00:00"
    },
    {
        "record_id": "sample_0021",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:14.999897+00:00"
    },
    {
        "record_id": "sample_0022",
        "prediction": 0,
        "timestamp": "2026-05-09T05:03:15.005760+00:00"
    },
    {
        "record_id": "sample_0023",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.009888+00:00"
    },
    {
        "record_id": "sample_0024",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.013844+00:00"
    },
    {
        "record_id": "sample_0025",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.007426+00:00"
    },
    {
        "record_id": "sample_0026",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.013892+00:00"
    },
    {
        "record_id": "sample_0027",
        "prediction": 1,
        "timestamp": "2026-05-09T05:03:15.019715+00:00"
    }
]
}

```




Figure 1: training phase

Short Answers

1. Describe your system end-to-end.

The system has two flows. In the **training flow**, an Airflow DAG (`ml_training_pipeline_v2`) loads the breast cancer dataset, splits it into train/test sets, trains a logistic regression classifier, evaluates it against a 0.94 accuracy threshold, and—if it passes—promotes the model by uploading versioned artifacts (`cancer.pkl`, `metrics.json`, `metadata.json`) and a `promoted_model.json` pointer to S3 (LocalStack). In the **inference flow**, a second Airflow DAG (`sqs_populate_inference_queue`) reads the same test split and sends one SQS message per record with a `record_id` and feature array. One or more Kubernetes consumer pods poll the SQS queue, load the promoted model from S3 on startup, run `model.predict()` on each message, write the result as `predictions/<record_id>.json` to S3, and delete the message only after a successful write.

2. Why is a queue used instead of direct API calls?

A queue decouples producers from consumers. The Airflow DAG can finish pushing all 113 test records instantly without waiting for any inference to complete. Consumers process at their own pace, and if they lag behind or crash, messages stay in the queue rather than being lost. A direct API call would require the caller to wait for a synchronous response, block on consumer availability, and implement its own retry logic for failures.




Figure 2: model info




Figure 3: prediction results

3. What happens if a consumer crashes mid-processing?

Each SQS message has a visibility timeout. When a consumer receives a message, SQS hides it from other consumers for that duration. If the consumer crashes before calling `delete_message`, the visibility timeout expires and the message reappears in the queue for another consumer to pick up. Because we only call `delete_message` after both the inference and the S3 write succeed, a mid-processing crash results in the message being retried—guaranteeing at-least-once delivery without losing any record.

4. Where are the bottlenecks in your system?

The primary bottleneck is the consumer’s sequential per-message processing loop: receive → predict → write to S3 → delete. S3 `put_object` adds network round-trip latency for every record. A secondary bottleneck is model loading: each consumer pod downloads the model from S3 once at startup, so cold-start time scales with model file size. The SQS queue itself is unlikely to be a bottleneck for this workload; the queue simply accumulates work that consumers drain at their own rate.

5. One improvement you would make for production.

Replace the ad-hoc `promoted_model.json` pointer with a proper model registry (e.g., MLflow). A registry provides a central API for versioning, stage transitions (staging → production), lineage tracking, and rollback. It removes the fragile convention of hand-crafting JSON pointers in S3 and makes it safe to run multiple model types and versions simultaneously without path collisions.